

# Programming with Data Bootcamp: Lecture 8

<http://dsg.csail.mit.edu/6.S079/>



Slides courtesy of Sam Madden  
/ Tim Kraska (6.S079)

## Key ideas:

Single-node Parallelism

Multi-node Parallelism

- Dask
- Spark
- Ray



# Parallelism Goal

- Make a job faster by running on multiple processors
- What do we mean by faster?


$$\text{speed up} = \frac{\text{old time}}{\text{new time}} \text{ on same problem, with } N \text{ times more hardware}$$

$$\text{scale up} = \frac{1x \text{ larger problem on } 1x \text{ hardware}}{Nx \text{ larger problem on } Nx \text{ hardware}}$$

- Not necessarily the same: smaller problem may be harder to parallelize

# Speedup Goal

- Linear?

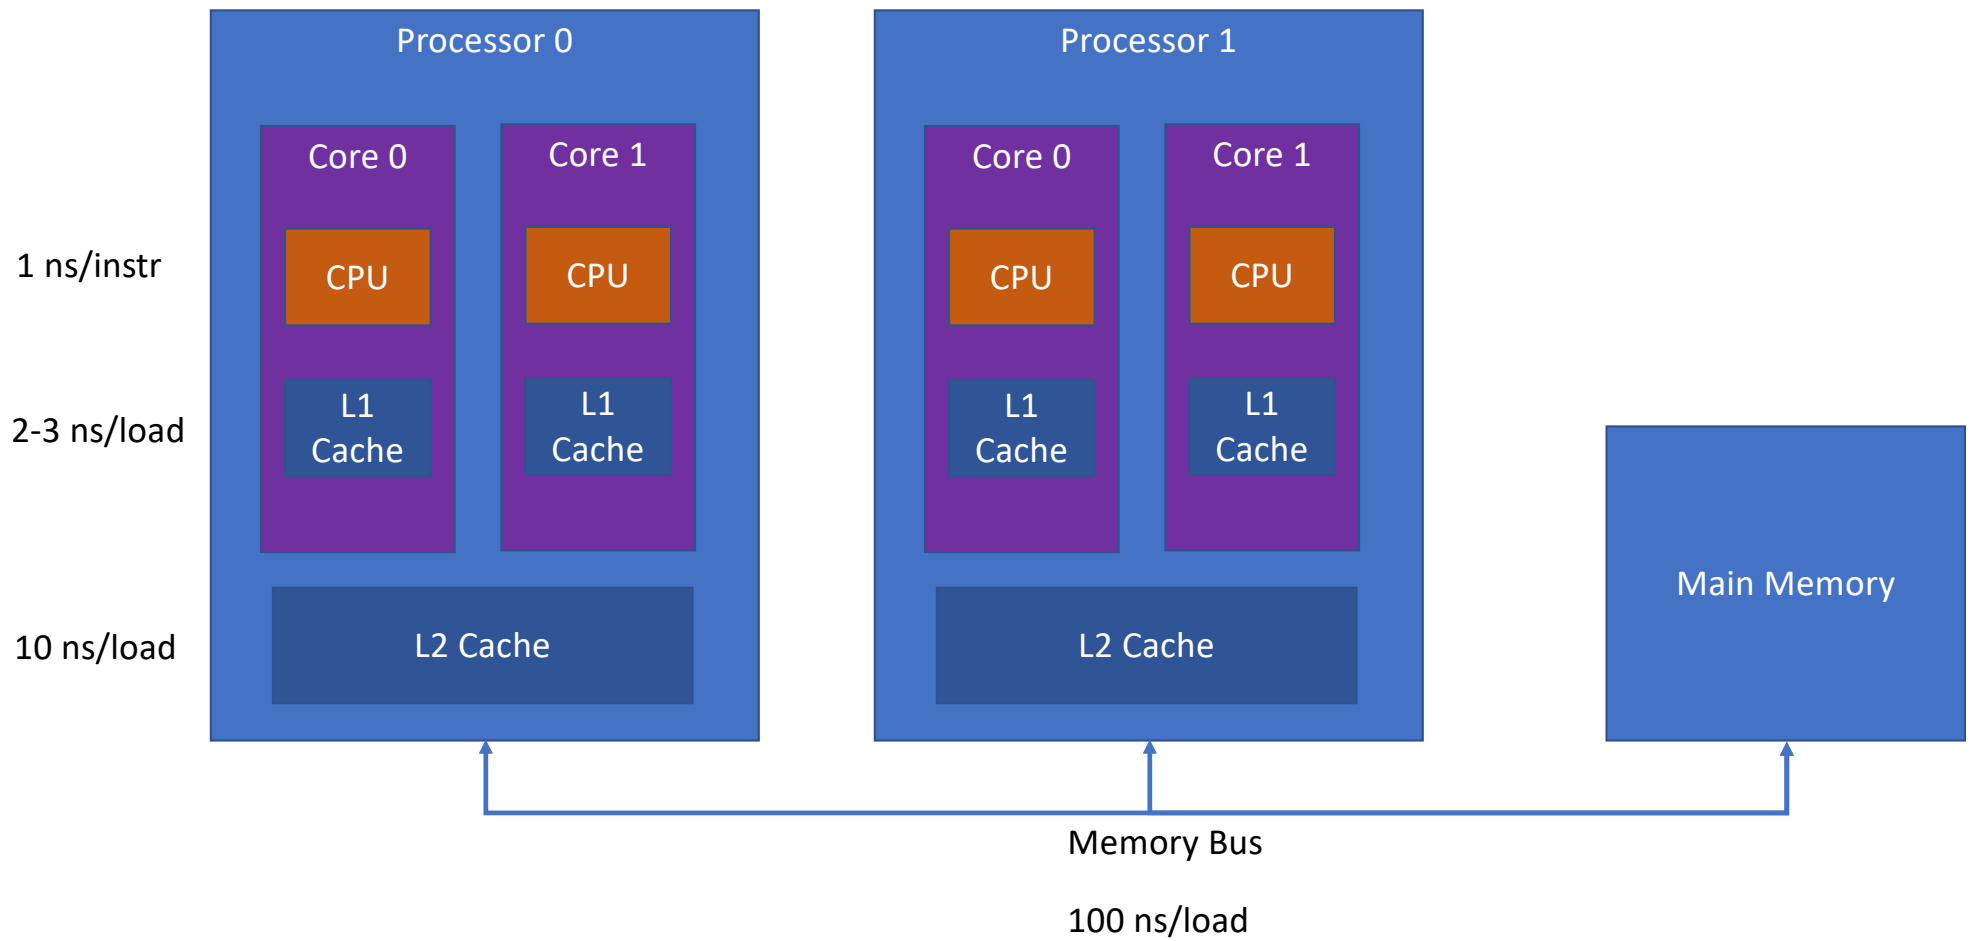


# Barriers to Linear Scaling


- Startup times
  - e.g., may take time to launch each parallel executor
- Interference
  - processors depend on some shared resource
  - E.g., input or output queue, or other data item
- Skew
  - workload not of equal size on each processor
- *Almost all workloads will stop scaling at some point!*
- What are some barriers in data science workloads?

# Properties of Parallelizable Workloads

- Provide linear speedup
- Usually can be decomposed into small units that can be executed independently
  - "embarrassingly parallel"
- As we will see, SQL-style operations generally provide this
- Some ML algorithms support it, but often tricky




Some machines may have 2 levels of cache per core



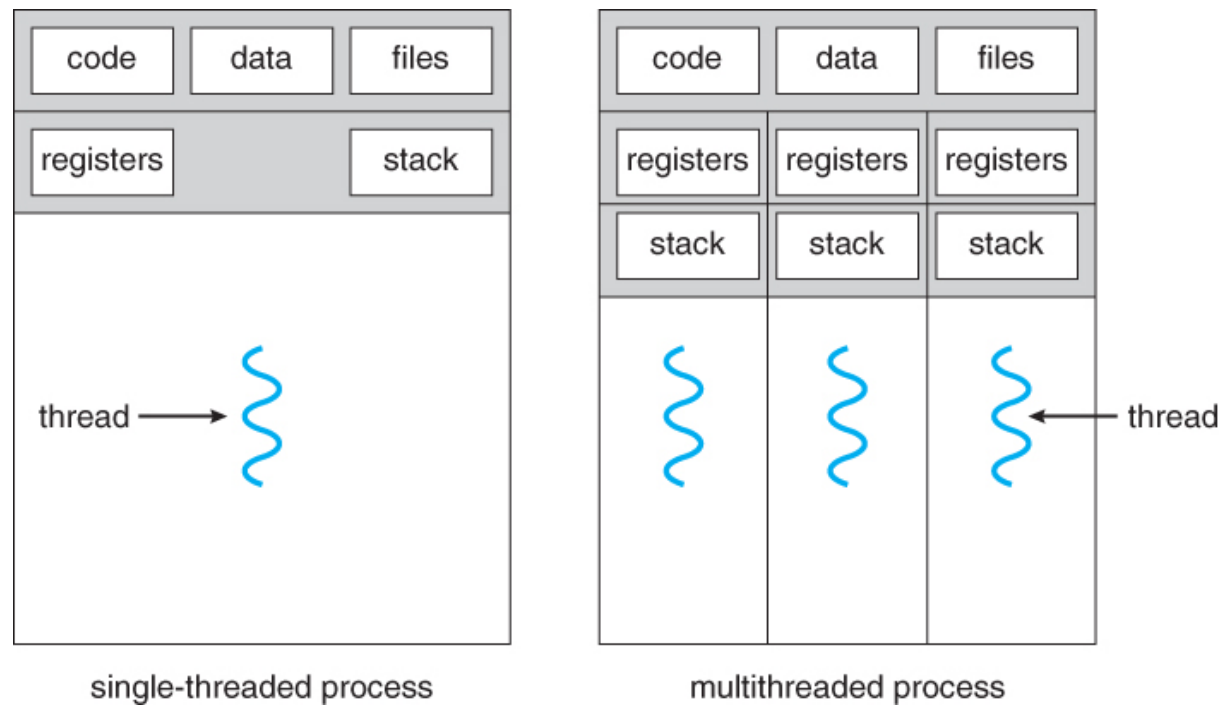
# Ping Test (Ethernet inside CSAIL)

- csail.mit.edu
  - 0.7 ms
- mit.edu
  - 14.0 ms
- harvard.edu
  - 7.0 ms
- berkeley.edu
  - 65.1 ms
- tsinghua.edu
  - 229.5 ms





# Threads vs Processes



[https://www.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/4\\_Threads.html](https://www.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/4_Threads.html)

# Python Threads API

```
import threading
```

```
t = threading.Thread(target=func_name, args=(a1,a2,...))  
t.start()    #start thread running – main thread continues  
t.join()     #wait for thread to finish
```

```
lock = threading.Lock()    #create a lock object  
lock.acquire() #acquire the lock; block if another thread has it  
lock.release() #release the lock
```

**Problem: Python Global Interpreter Lock (GIL)**  
**Only one thread can be executing python code at once**

# Python Multiprocessing API

```
import multiprocessing
```

```
p = multiprocessing.Process(target=func_name, args=(a1,a2,...))
```

```
p.start()    #start thread running – main thread continues
```

```
p.join()    #wait for thread to finish
```


```
lock = multiprocessing.Lock()    #create a lock object
```

```
lock.acquire() #acquire the lock; block if another thread has it
```

```
lock.release() #release the lock
```

# Parallel Aggregation

Task: compute average age across all people



```
{"age": 30, "name": ["Michal", "Sharpe"],  
"occupation": "Archivist", "telephone":  
"285.290.9033", "address": {"address":  
"458 Girard Plantation", "city":  
"Wentzville"}, "credit-card": {"number":  
"5384 0033 6904 0042", "expiration-date":  
"06/23"}}
```

# Parallel Aggregation Implementation

- Use multiprocessing, not threading
- Main thread creates a work queue

```
q = multiprocessing.Queue()
```

- Puts work on it, as pointers to files

```
q.put(file1); q.put(file2)
```

- Starts threads, passing them the work queue, as well as a result queue
- Threads pull from queue in a loop:


```
while True:  
    f = q.get(block=False)  
    process(f)
```

- Threads compute running sum and average
- Once complete, threads put their running sum and average on the result queue:

```
out_q.put((age_sum, age_cnt))
```

- Main thread blocks on result queue to read a result from each worker:

```
for p in procs:  
    (p_sum, p_count) = out_q.get()
```




# Clicker

Why didn't this program speed up beyond 8 processes? Choose all that apply

- a) Not enough memory
- b) Not enough processors
- c) Startup overheads of launching processes
- d) Too much coordination between processes

<https://clicker.mit.edu/6.S079>

# Break





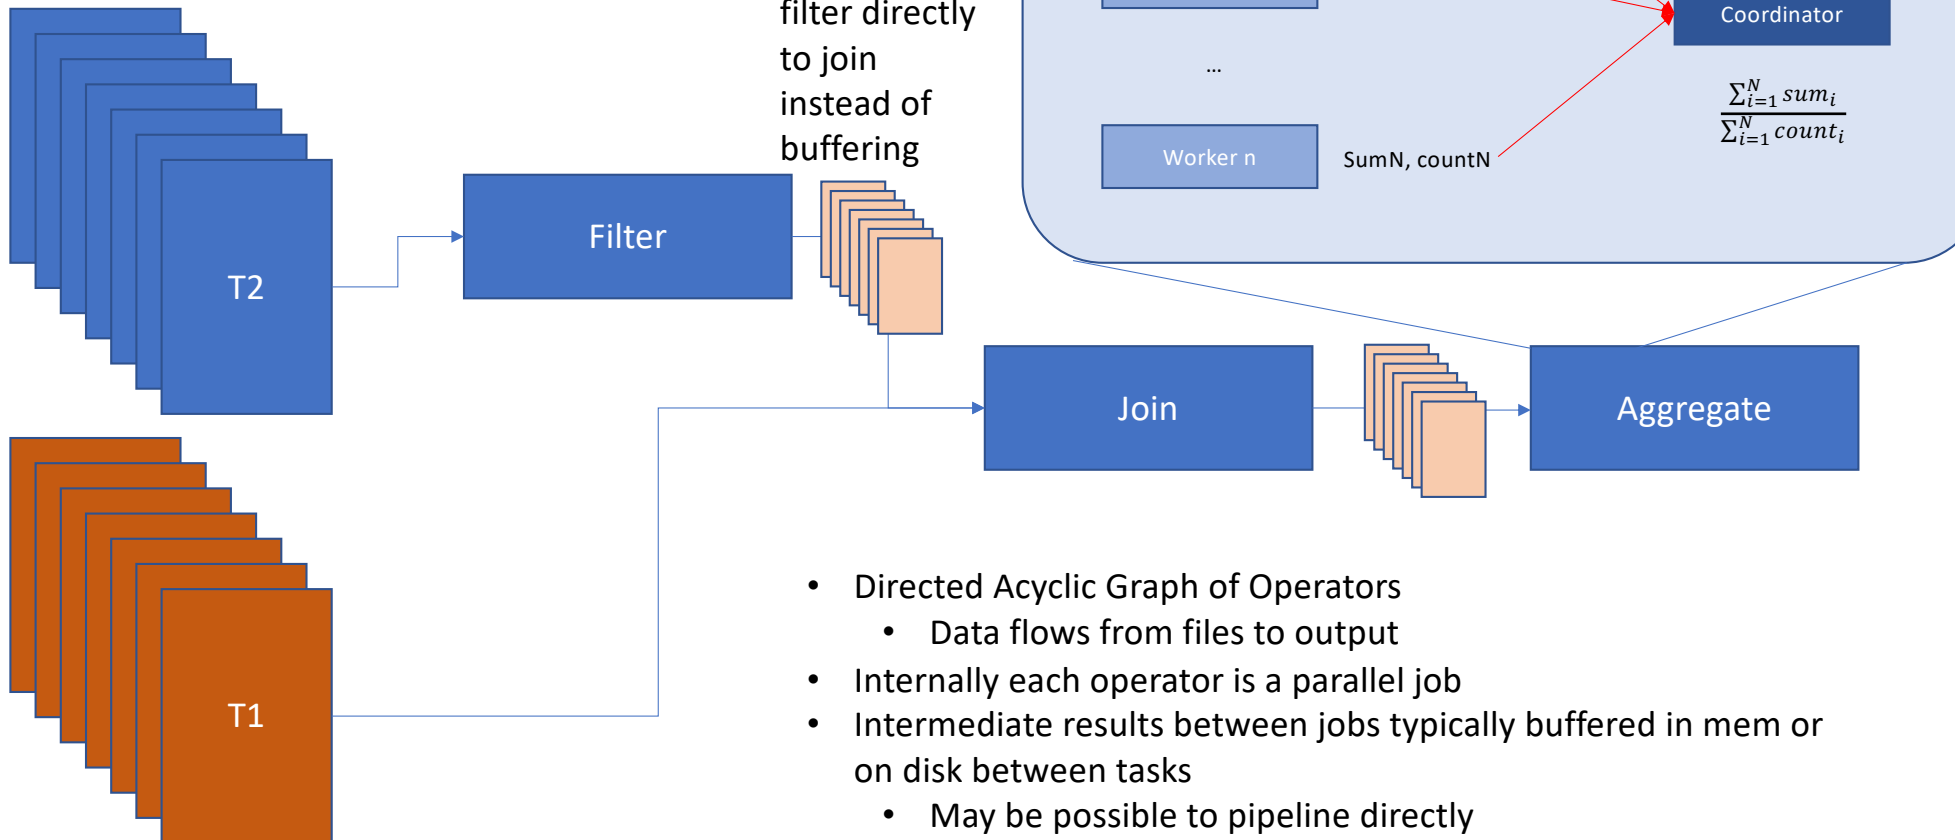
# Parallelism Approach

Split given data set into  $N$  partitions

Use  $M$  processors to process this data in parallel

We will need to come up with parallel implementations of common operators

# Parallel Dataflow Example



# Parallel Dataflow Operations


- Filter
- Project
- Element-wise or row-wise transform
- Join
  - Repartition vs broadcast
- Aggregate
- Sort
- Train an ML model
- Arbitrary python "UDF"s

*Which of these are easy to parallelize?*

# Partitioning Strategies

- Random / Round Robin
  - Evenly distributes data (no skew)
  - Requires us to repartition for joins
- Range partitioning
  - Allows us to perform joins/merges without repartitioning, when tables are partitioned on join attributes
  - Subject to skew
- Hash partitioning
  - Allows us to perform joins/merges without repartitioning, when tables are partitioned on join attributes
  - Only subject to skew when there are many duplicate values

# Round Robin Partitioning




## Advantages:

Each partition has the same number of records

## Disadvantage:

No ability to push down predicates to filter out some partitions

# Range Partitioning



## Advantages:

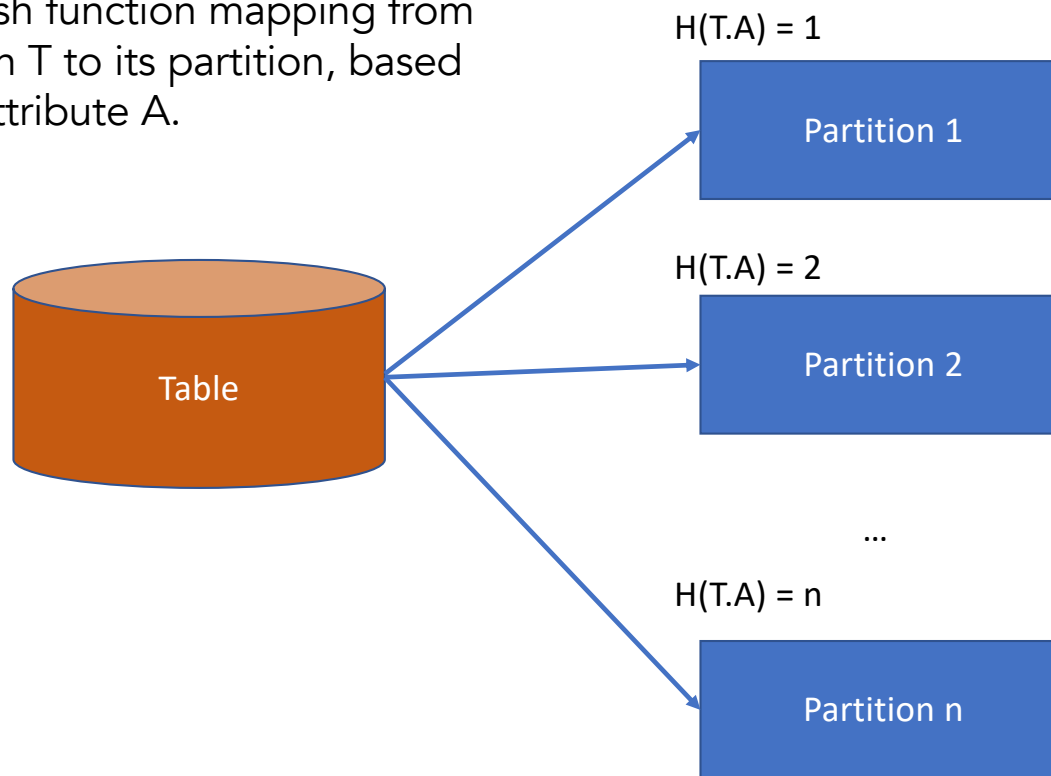
Easy to push down predicates (on partitioning attribute)

## Disadvantage:

Difficult to ensure equal sized partitions, particularly in the face of inserts and skewed data

# Hash Partitioning

$H(T.A)$  is a hash function mapping from each record in  $T$  to its partition, based on value of attribute  $A$ .



## Advantages:

Each partition has about the same number of records, unless one value is very frequent

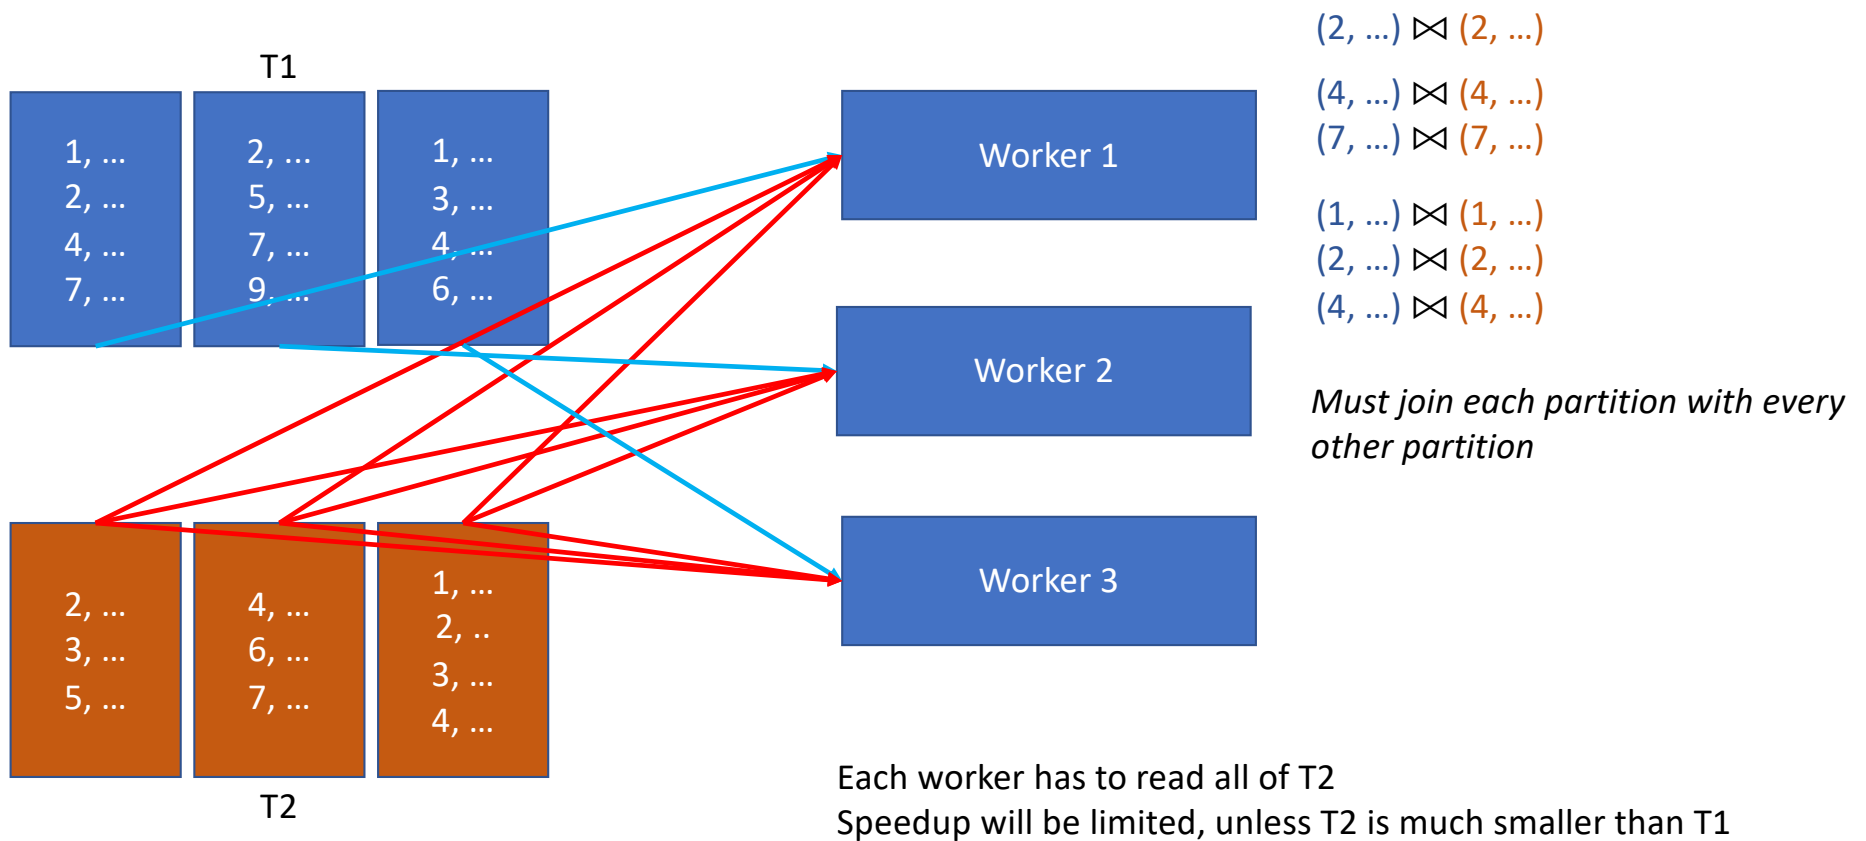
Possible to push down equality predicates on partitioning attribute

## Disadvantages:

Can't push down range predicates

# Parallel Join – Random Partitioning Naïve Algo

(1, ...) indicates value of join attribute






# Parallel Join – Prepartitioned

(1, ...) indicates value of join attribute

*Only need to join partitions that match*



(1, ...) ⋈ (1, ...)  
(1, ...) ⋈ (1, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)

*This is what our Postgres example showed*

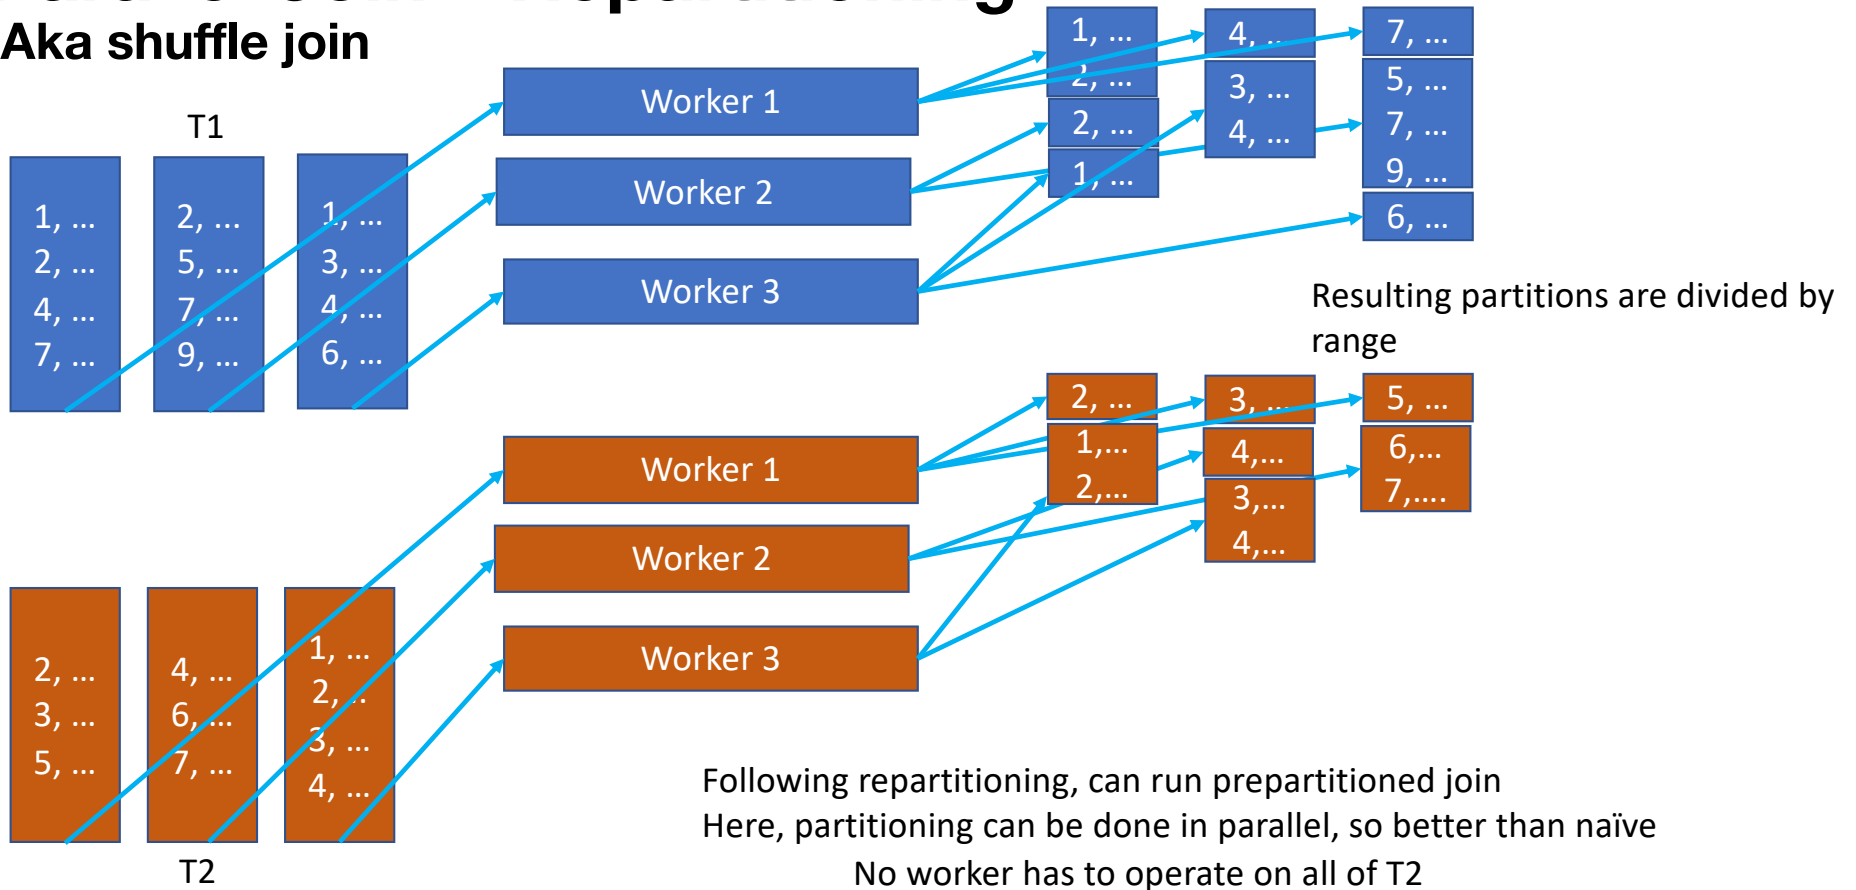
Better speedup, only works if data is properly prepartitioned

Should be 3x faster than single node join

Skew problem (hashing may help)

# Parallel Join – Repartitioning

Aka shuffle join



# Dask

<https://dask.org>




- General purpose python parallel / distributed computation framework
- Includes parallel implementation of Pandas dataframes
- Usually straightforward to translate a pandas program into a parallel implementation
  - Just use `dask.dataframe` instead of `pandas.dataframe`
  - Have to specify a parallel configuration to run on, via `Client()` object
    - Can be a local machine or distributed cluster
- Also has support for other types of parallelism, e.g., `dask.bag` class that allows parallel operation on collections of python objects

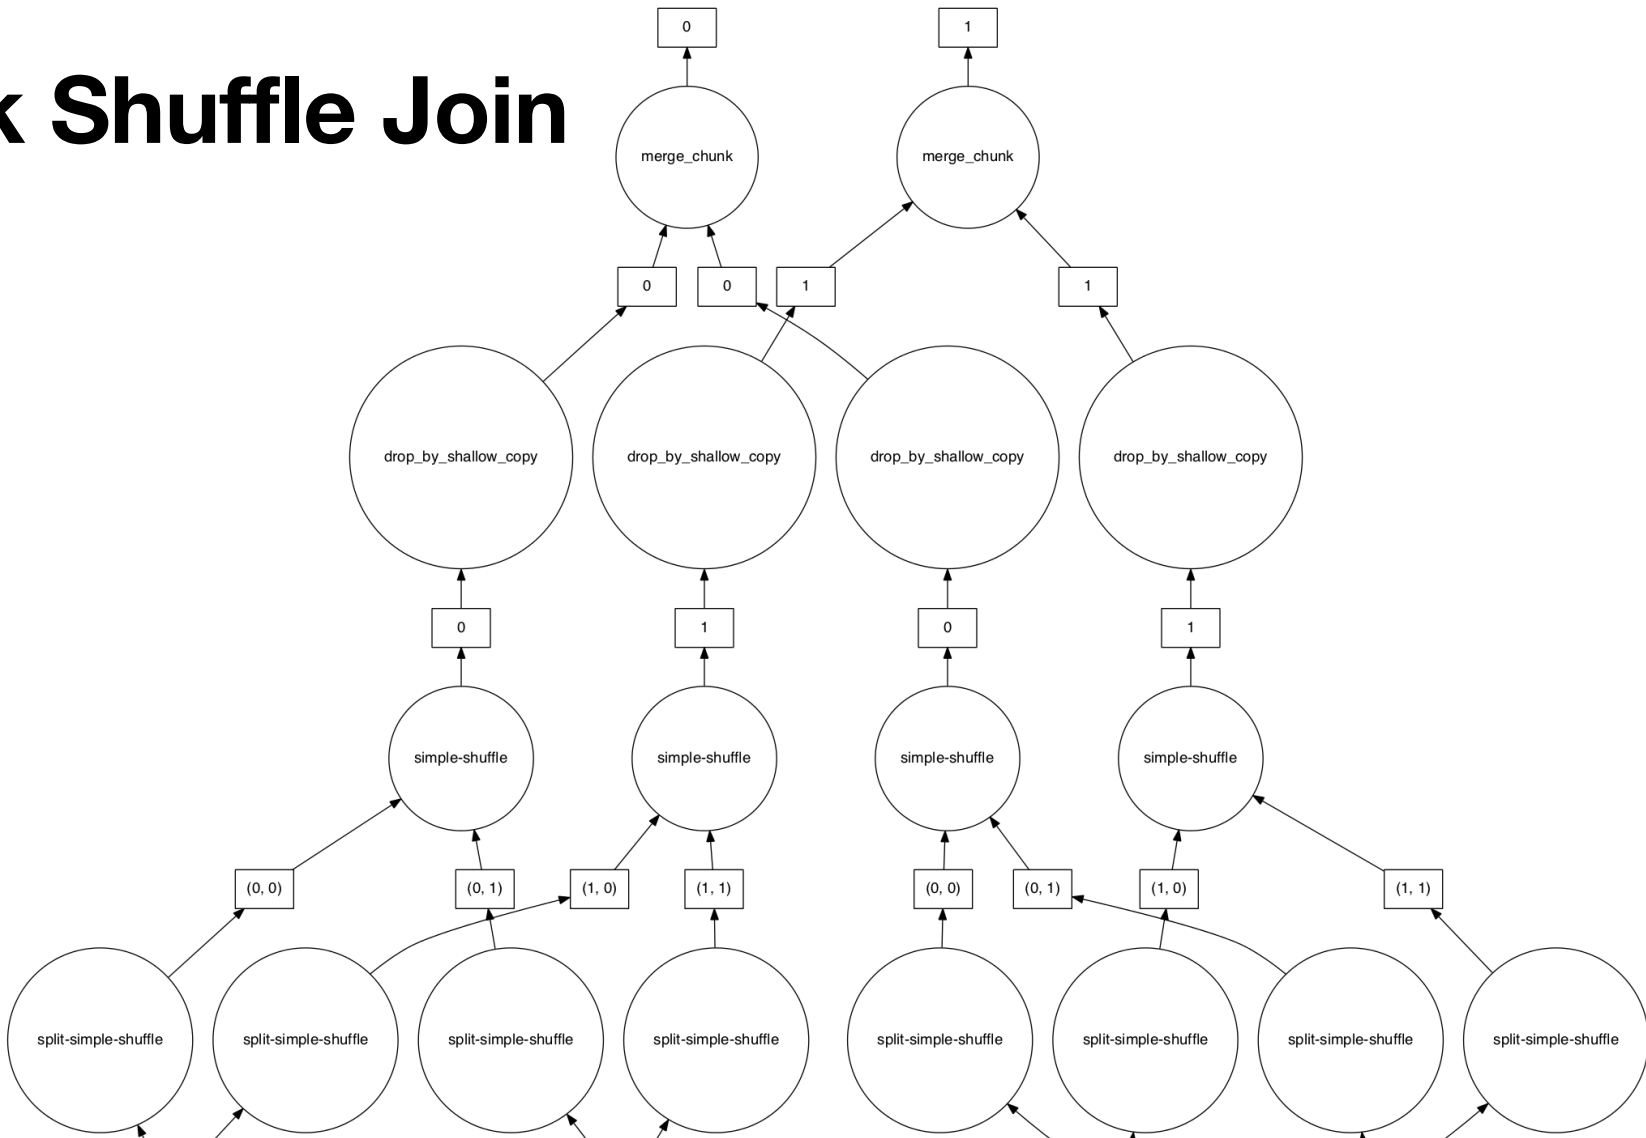
# Large Join Demo

- Changing number of nodes
- Changing join algorithm

# Dask Partitioned Join



# Dask Shuffle Join

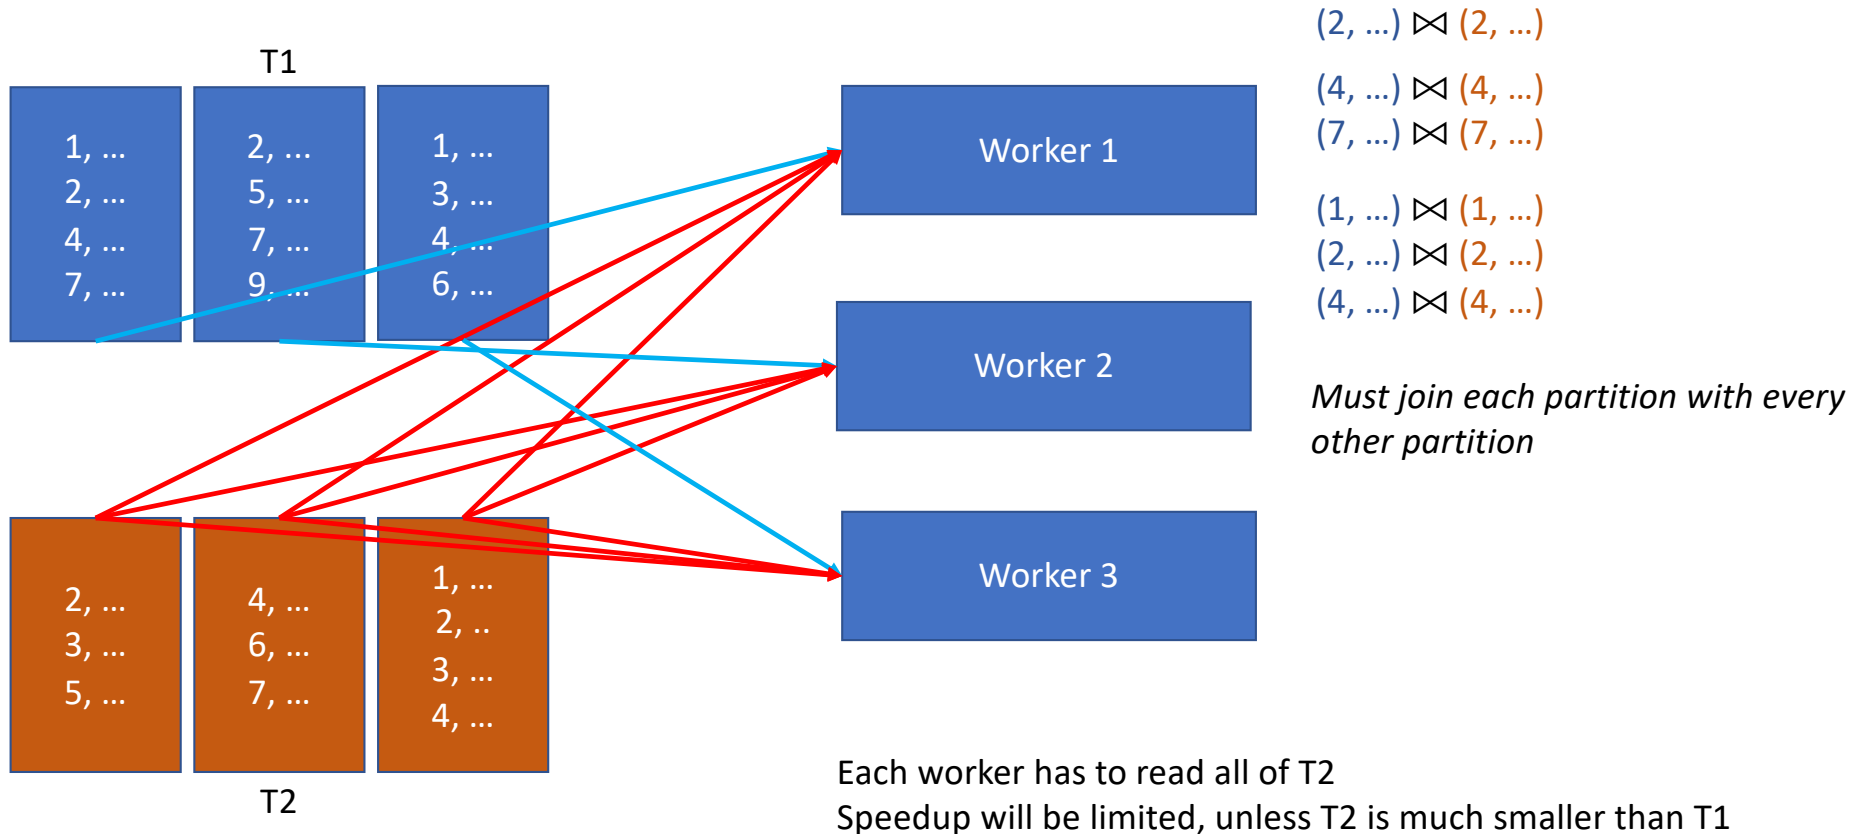


# Many alternatives

- MapReduce / Hadoop
  - Rewrite you program as collection of parallel map() and reduce() jobs
  - Hard to do, slow()
- Spark
  - Popular library -- similar to dask, more focused on large scale distributed
  - Includes parallel implementations of ML and other operations
  - Difficult to use

# Parallel Join – Random Partitioning Naïve Algo

(1, ...) indicates value of join attribute






# Parallel Join – Prepartitioned

(1, ...) indicates value of join attribute

*Only need to join partitions that match*



(1, ...) ⋈ (1, ...)  
(1, ...) ⋈ (1, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)  
(2, ...) ⋈ (2, ...)

*This is what our Postgres example showed*

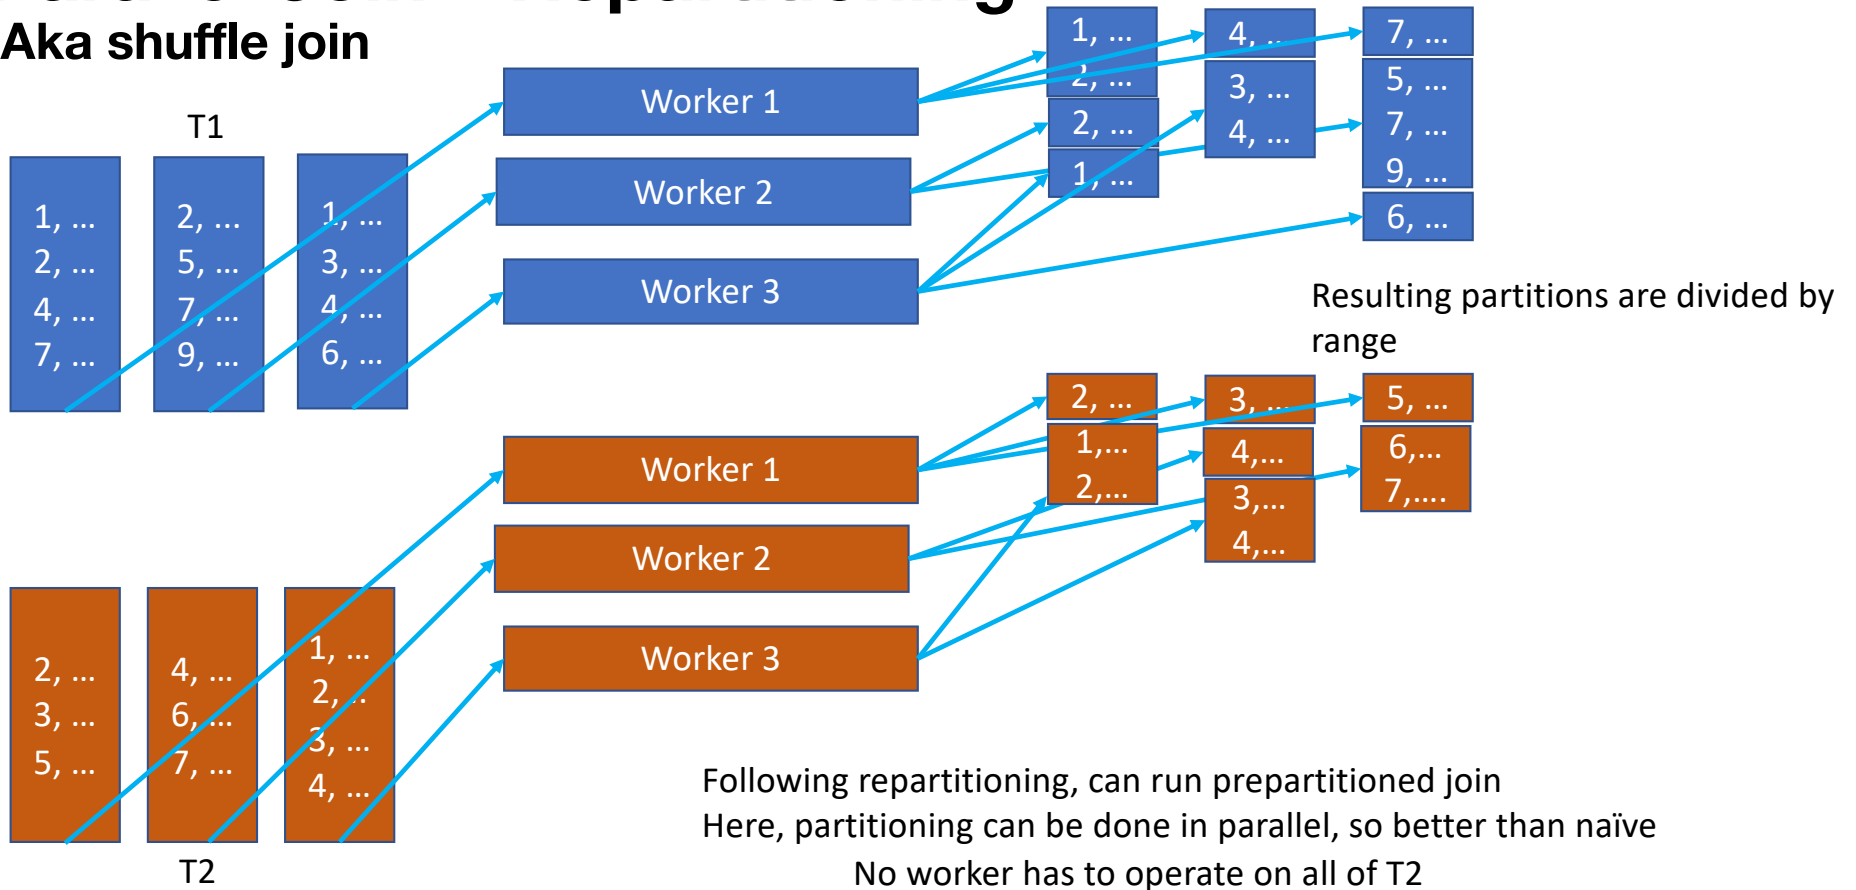
Better speedup, only works if data is properly prepartitioned

Should be 3x faster than single node join

Skew problem (hashing may help)

# Parallel Join – Repartitioning

Aka shuffle join



# Recap: Large Join In Dask

```
client = Client(n_workers=8, threads_per_worker=1, memory_limit='16GB')

header = "CMTE_ID,AMNDT_IND,RPT_TP,TRANSACTION_PGI,IMAGE_NUM,TRANSACTION_TP ..."
PATH = "indiv20/by_date/itcont_2020_20010425_20190425.txt"
PATH2 = "indiv20/by_date/itcont_2020_20190426_20190628.txt"

df = dask.dataframe.read_csv(PATH, low_memory=False, delimiter='|', header=None ...)
df2 = dask.dataframe.read_csv(PATH2, low_memory=False, delimiter='|', header=None ...)
df = df.dropna(subset=['NAME']).drop_duplicates(subset=['NAME'])
df2 = df2.dropna(subset=['NAME']).drop_duplicates(subset=['NAME'])

# make 3 copies
df = df.append(df)
df = df.append(df)
df = df.append(df)

df2 = df2.append(df2)
df2 = df2.append(df2)
df2 = df2.append(df2)

ans = df.merge(df2, on='NAME').count()

ans = ans.compute()      Execution is deferred until compute is called

print(f"found {ans} matches")
```

# Dask Distributed

*"Distributed" = multiple machine*

*"Parallel" = multiple processors on same machine*

- Demo on Amazon
  - Much slower than laptop, t3.large machines (8GB RAM, 2x vCPU ~30% performance / CPU)
- Single local executor: 174.3 s
- Single distributed worker: 200.5
- Three distributed workers: 78.5 s (2.2x/2.6 speedup)

# Subgraph Caching via “Persist”

- Can “persist” a subresult to cause it to be stored in memory
- Avoids recomputing

```
n1 = df.loc[:, ["NAME"]].persist()
n2 = df2.loc[:, ["NAME"]].persist()

#will compute the count and persist n1 and n2
ans = n1.merge(n2, on='NAME').count()
print(ans.compute())

#will reuse previously persisted result
ans2 = n1.merge(n2, on='NAME').max()
print(ans2.compute())
```

# Fault Tolerance Model

- Retries tasks that fail
- Resilient to the failure of any one worker
- Demo

# Spark

- Distributed / parallel data processing system
- pyspark.sql engine very similar to dask in functionality
  - Slightly different API
  - Other pandas-on-spark projects, e.g., koalas provide pandas API compatibility

# Example

Demo!

```
spark = SparkSession.builder.appName("SimpleApp").getOrCreate()

path = "indiv20/by_date/itcont_2020_20010425_20190425.txt"
path2 = "indiv20/by_date/itcont_2020_20190426_20190628.txt"
header = "CMTE_ID,AMNDT_IND,RPT_TP,TRANSACTION_PGI,IMAGE_NUM,TRANSACTION_TP, ..."

df_spark = spark.read.csv(path, sep='|', header=False)
df_spark = df_spark.toDF(*header)
df_spark = df_spark.dropna(subset=["NAME"]).dropDuplicates(subset=["NAME"])
df_spark = df_spark.union(df_spark)
df_spark = df_spark.union(df_spark)
df_spark = df_spark.union(df_spark)

df_spark2 = spark.read.csv(path2, sep='|', header=False)
df_spark2 = df_spark2.toDF(*header)
df_spark2 = df_spark2.dropna(subset=["NAME"]).dropDuplicates(subset=["NAME"])
df_spark2 = df_spark2.union(df_spark2)
df_spark2 = df_spark2.union(df_spark2)
df_spark2 = df_spark2.union(df_spark2)

ans = df_spark.join(df_spark2, on='NAME').count()
print(ans)
```

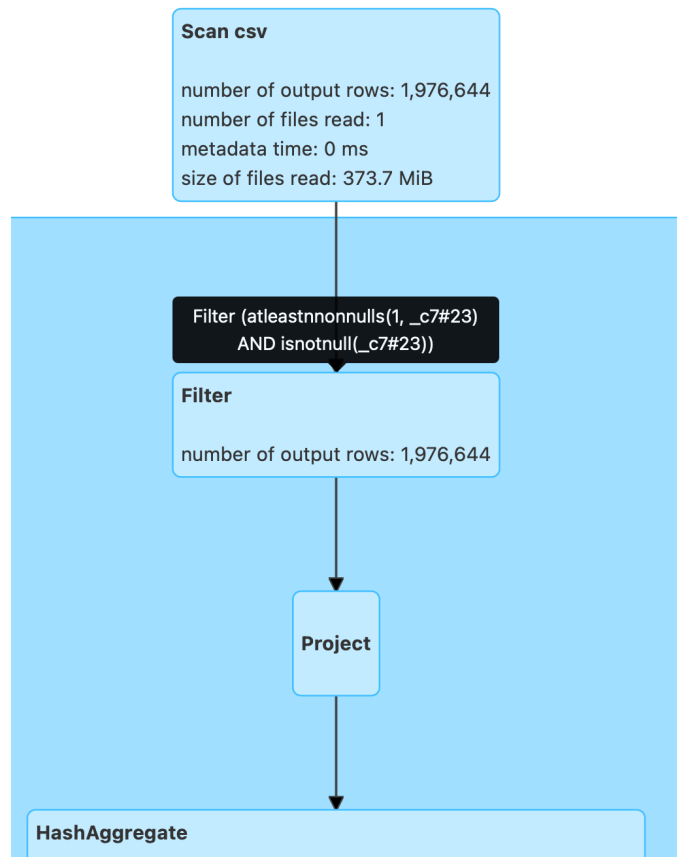
*This is a way to run spark locally;  
most people run a cluster of machines  
and submit jobs, like the dask  
distributed demo before*



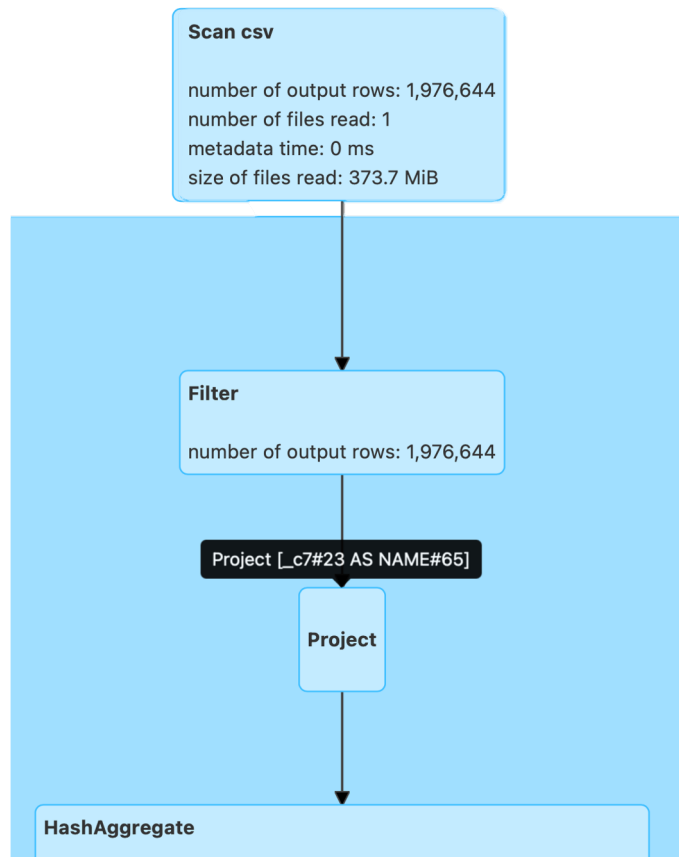
# Spark Under the Hood

- Compiles to Java/Scala
  - Makes understand what tasks are doing and debugging messages somewhat confusing
- Query optimizer much smarter than Dask
  - Projection push down
  - Pre-aggregation

# Projection Push Down

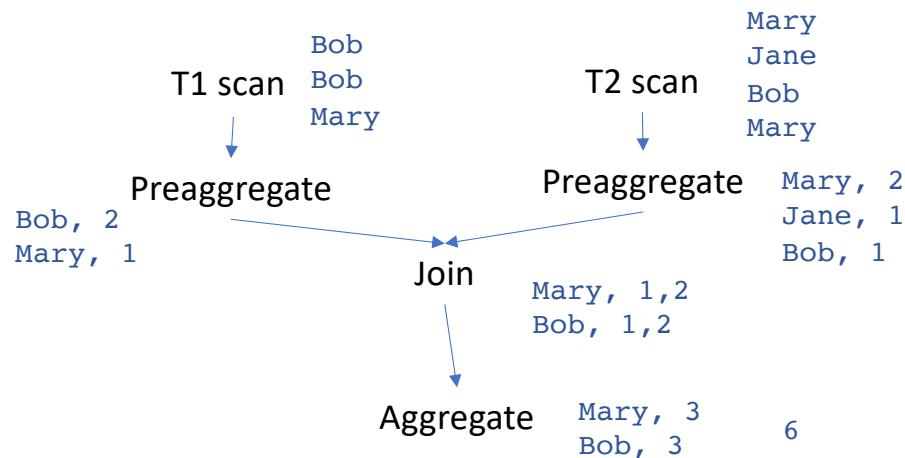


# Projection Push Down



# Preaggregation

- Goal: count the number satisfying records in the join
- Idea: count records in each table before the join
- Join {record, count} pairs from tables to compute final join
- Eliminates the number of records that need to join



*In spark, preaggregate, join and aggregate can all be done massively in parallel*

# Spark vs Dask

- Dask is much smaller, more pythonic
- Spark generally performs better
  - More optimized for very large datasets on S3 / cloud storage
  - Dask lacks query optimization
- Spark is harder to use and debug
  - Compilation down to Java makes it hard to understand what is happening, sometimes
- Many other packages in spark, including
  - SparkML
  - Spark Streaming
  - A variety of data lake / storage tools

# Summary

- Dask and Spark both support parallel and distributed computation over data
  - Both scale from a few processors to hundreds of machines
- Dask is good for parallelizing pandas/numpy code
- Spark more sophisticated, less tied to python ecosystem